

AI Governance Framework

Office of the CTO | Version 1.0 | Approved by Board January 22, 2026

INTERNAL — All Employees

1. AI Governance Structure and Authority

The AI Governance Framework was approved by the Board on January 22, 2026 and applies to all production AI/ML models deployed by TechNova on the Nova Platform, whether developed in-house or sourced from third parties. It does not apply to research prototypes that are clearly labelled non-production.

Governance is exercised by the AI Ethics Committee (AIEC), constituted of five voting members: the Chief Technology Officer (chair), the Chief Information Security Officer, the Chief Revenue Officer (representing customer impact), one independent director (currently Dr. Vikram Reddy), and the General Counsel. The AIEC meets monthly; meeting minutes are board-visible. Quorum requires three voting members including the chair or the General Counsel.

The AIEC has exclusive authority to: (i) approve any new AI model's promotion to production, (ii) approve material changes (defined below) to a production model, (iii) suspend any production model in response to a fairness, safety, or compliance concern, and (iv) approve customer-facing communication that describes a model's behaviour.

2. Model Lifecycle and Material Change Definition

Every production AI model passes through five gates: (G1) Concept Approval, (G2) Training Data Approval, (G3) Bias and Fairness Testing, (G4) Pre-Production Risk Review, (G5) Post-Deployment Monitoring Plan. Gates G1 through G4 require AIEC sign-off; G5 is operational and managed by the SRE-AI team.

A material change requiring re-approval is any of the following: (i) any change to the training dataset that adds or removes more than 5% of records; (ii) any change to the loss function or model architecture; (iii) any change to a feature input that materially affects a protected attribute; (iv) any deployment to a new geographic market where the model has not been previously validated; (v) any change in inference latency exceeding 25% of the prior baseline.

Non-material changes (e.g. infrastructure updates, batch-size tuning that does not affect outputs) may be approved by the SRE-AI team alone without an AIEC vote, with the change captured in the model card.

3. Bias, Fairness, and Safety Testing

Every production model is tested for disparate impact across the following protected attributes: gender, age band, region (with India-specific tests for regional bias), industry (with specific BFSI and

Healthcare bias panels because of regulatory sensitivity), and customer-tier. The current bias-testing suite uses Aequitas, Fairlearn, and an in-house regulatory checklist.

Acceptance thresholds: (i) maximum disparity of 10 percentage points between the highest-performing and lowest-performing protected group on the model's primary metric (typically AUC or F1); (ii) for credit and risk-scoring use cases (which TechNova does not currently sell but may in future), the threshold tightens to 4 percentage points; (iii) any disparity exceeding 5 percentage points triggers an automatic AIEC review.

Safety testing for generative AI features specifically includes: prompt injection resistance, data exfiltration via prompts (no customer data should leave the model context), refusal-rate calibration, and citation-faithfulness for the AI Co-pilot feature. Generative outputs that cite sources must achieve 95% faithfulness in QA samples; AI Co-pilot is currently at 96.2%.

4. Model Risk Tiers and Approval Authority

TechNova classifies AI models into three risk tiers, each with distinct approval requirements.

Tier A (High Risk) models include any model whose output directly drives an automated decision affecting customer money, customer access, or regulated outcomes. These require unanimous AIEC approval at G1, G2, G3, and G4 gates. Tier A models also require an external audit by a recognised AI assurance firm before launch; no Tier A models are currently in production at TechNova.

Tier B (Medium Risk) models include recommendation systems that materially influence customer decisions but where a human reviews the output, and forecasting models used in operational planning. The AI Co-pilot, Predictive Forecasting (ML), and AI-Powered Anomaly Detection features fall in this tier. They require AIEC majority approval (3 of 5 votes) at all gates and quarterly post-deployment monitoring reviews.

Tier C (Low Risk) models include narrow utility models such as ranking, sorting, and search-relevance scoring where output is non-binding and easily verifiable by the user. Tier C models can be approved by the CTO and CISO jointly without full AIEC voting.

The current Tier B inventory: TechNova-Classify (BERT for document classification), TechNova-Extract (custom NER for data extraction), TechNova-Assist (RAG-based customer support powered by Qwen-2.5-72B with LoRA adapters), Predictive Forecasting (ML), AI-Powered Anomaly Detection, and AI Co-pilot (Beta).

5. Production Monitoring, Drift, and Deprecation

Every production model has a continuously-monitored model card containing: model version, training-data summary, bias-test results, performance benchmarks, deployment regions, last AIEC review date, and known limitations. Model cards are accessible to Customer Success teams and on request to enterprise customers under NDA.

Drift monitoring uses Evidently AI with thresholds tuned per model. Data drift exceeding 10% on the primary feature distribution triggers a yellow alert; exceeding 20% triggers a red alert and a 5-day

window for remediation. Concept drift (degradation of primary metric) follows the same thresholds.

A model deprecation is initiated when (i) the primary metric degrades by more than 15% from launch baseline and cannot be remediated within 30 days, (ii) a successor model with superior fairness or accuracy is launched, or (iii) AIEC determines the use case is no longer aligned with TechNova's policies. Deprecation provides 90-day notice to affected customers and 180-day total runway before final retirement.

The current FY2026-27 plan allocates 38% of the INR 485 crore total product investment budget to AI/ML capabilities (= INR 184.30 crores), with infrastructure including the planned self-hosted RAG stack on 4x A100 nodes with customer-specific Qdrant vector indices.

6. Customer-Facing AI Disclosure Requirements

TechNova's contractual commitments to customers regarding AI feature behaviour are governed by this section.

Disclosure: every customer using a Tier-A or Tier-B AI feature receives a model card excerpt at activation, summarising what the model does, what data it uses, what its known limitations are, and what disparate-impact testing has been done. The customer's MSA Data Processing Addendum (DPA) governs how that customer's data may be used in model training; by default, customer data is NOT used to train shared models without explicit opt-in.

Customer rights: customers may request that AI predictions affecting them be reviewed by a human within 5 business days. Customers may request the deactivation of any AI feature on their tenant within 24 hours. Customers may request a copy of bias-testing results under NDA.

Audit support: TechNova will support customer auditors examining AI model usage on their data with up to 8 hours of engineer time per quarter, included in the standard MSA. Additional audit support is billable.

Regulatory disclosures: in jurisdictions with AI-specific regulations (currently the EU AI Act and emerging India AI Act draft), TechNova maintains a regulatory-mapping document, last refreshed in March 2026, that maps each Tier-A/B model to applicable regulatory categories.